

Introduction to Kubernetes

The Hard Way

July 1, 2022

Table of Contents

1. Background
2. Setup
3. Concept
4. Operator
5. Summay

Background

Why Kubernetes?

Container's **Unique** Advantages


Lightweight footprint

Minimal overhead

Portability across machines

More consistent operation

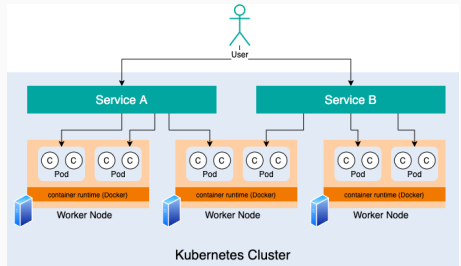
Acceptable Isolation



Why Kubernetes?

DevOps challenges for multiple containers

Management
Metrics
Health checks
Networking
Scheduling
Scaling
Deployment
Rollbacks
Zero-downtime



Why Kubernetes?

Meet **Kubernetes**

Now the most popular container cluster

Manage applications, not machines

Open source, written in Go



Setup

VirtualBox

Oracle Virtual Machine

Vagrant

build and maintain portable virtual machine
through command line

Goal

one master node + 2 worker nodes

=> kubernetes cluster

vagrant up

Vagrantfile

```
MASTER_IP      = "192.168.56.10"
NODE_01_IP      = "192.168.56.11"
NODE_02_IP      = "192.168.56.12"
Vagrant.configure("2") do |config|
  config.vm.box = "geerlingguy/ubuntu2004"
  config.vm.box_version = "1.0.3"
  boxes = [
    { :name => "kube-master",  :ip => MASTER_IP,  :cpus
      => 2, :memory => 4096 },
    { :name => "kube-node-01", :ip => NODE_01_IP, :cpus
      => 2, :memory => 4096 },
    { :name => "kube-node-02", :ip => NODE_02_IP, :cpus
      => 2, :memory => 4096 },
  ]
  // ...
```

Vagrantfile

```
boxes.each do |opts|
  config.vm.define opts[:name] do |box|
    box.vm.hostname = opts[:name]
    box.vm.network :private_network, ip: opts[:ip]
    box.vm.provider "virtualbox" do |vb|
      vb.cpus = opts[:cpus] vb.memory = opts[:memory]
    end
    box.vm.provision "shell", path:"./install.sh"
    if box.vm.hostname == "kube-master" then
      box.vm.provision "shell", path:"./cfg-master.sh"
    else
      box.vm.provision "shell", path:"./cfg-workers.sh"
    end
  end
end
end
```

Guest - master node

cfg-master.sh

```
// ...
initialize_master_node () {
sudo systemctl enable kubelet
sudo kubeadm init
    --apiserver-advertise-address=$master_node
    --pod-network-cidr=$pod_network_cidr
    --image-repository
registry.aliyuncs.com/google_containers
    --ignore-preflight-errors=NumCPU
}

create_join_command () {
kubeadm token create --print-join-command | tee
    /vagrant/join_command.sh
chmod +x /vagrant/join_command.sh
}
```

cfg-workers.sh

```
#!/bin/bash -e

get_join_command () {
sudo /vagrant/join_command.sh
}

get_join_command
```

vagrant ssh kube-master

on master node

```
vagrant@kube-master:/vagrant$ kubectl get node
```


NAME	STATUS	ROLES	AGE	VERSION
kube-master	Ready	control-plane,master	8d	v1.20.11
kube-node-01	Ready	<none>	8d	v1.20.11
kube-node-02	Ready	<none>	8d	v1.20.11

Concept

Node


Basic Unit that Forms a Cluster

- **Master/Slave** Architecture
- Master node looks after
 - scheduling and scaling applications
 - maintaining the state of the cluster
 - implementing updates
- Worker nodes are responsible for
 - running the containers
 - doing any work assigned by the master node



Fundamental Kubernetes Work Unit

- Group of one or more containers
- Live and die together
- Support co-located processes
- Share
 - IP
 - Labels
 - Volumes



kubectl apply -f flask-pod.yaml

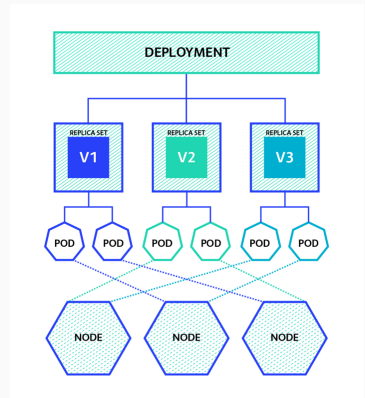
flask-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: flask-pod
  labels:
    app: flask-helloworld
spec:
  containers:
    - name: flask
      image: digitalocean/flask-helloworld:latest
      ports:
        - containerPort: 5000
```

Deployment

Manage Multiple Pods

- Try to maintain the desired state
- Creates new resources or replaces the existing resources when necessary
 - Life Cycle
 - Scheduler
 - Scaling



Deployment

kubectl apply -f flask-deployment.yaml

flask-deployment.yaml


```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: flask-dep
  labels:
    app: flask-helloworld
spec:
  replicas: 2
  selector:
    matchLabels:
      app: flask-helloworld
// ...
```

flask-deployment.yaml

```
// ...  
template:  
  metadata:  
    labels:  
      app: flask-helloworld  
  spec:  
    containers:  
      - name: flask  
        image: digitalocean/flask-helloworld  
        ports:  
          - containerPort: 5000
```

Abstraction of Logical Pods

- Routes traffic across a set of Pods
- Allow pods to die and replicate without impacting applications
 - ClusterIP
 - NodePort
 - LoadBalancer



```
kubectl apply -f flask-service.yaml
```

flask-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: flask-svc
  labels:
    app: flask-helloworld
spec:
  type: LoadBalancer
  ports:
    - port: 80
      targetPort: 5000
      protocol: TCP
  selector:
    app: flask-helloworld
```

Operator

Resource Object

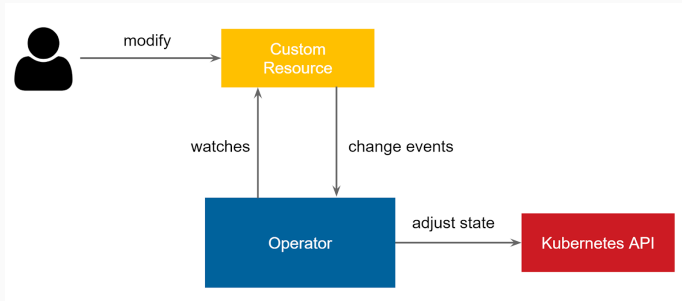
A resource is essentially a **collection of similar objects**.

- Built-in resources we already know
 - Pod
 - Deployment
 - Service
- Custom resources
 - A way to extend it and customize kubernetes resources.



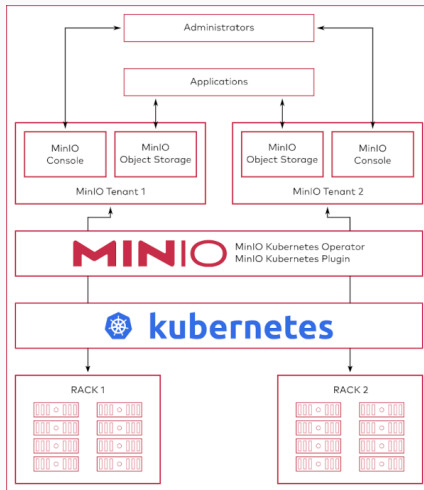
Controller

Each controller tries to move the current cluster state closer to the desired state.



What's Operator?

operator = custom resource definition + controller



Deploy MinIO on Kubernetes

Without Specific Operator

minio-pv.yaml

```
apiVersion: v1
kind: PersistentVolume
metadata:
  labels:
    app: minio
    release: minio
  name: minio
  namespace: default
spec:
  accessModes: - ReadWriteOnce
  capacity: storage: 10Gi
  volumeMode: Filesystem
  hostPath:
    path: /mnt/minio
```

Deploy MinIO on Kubernetes - Without Specific Operator

minio-pvc.yaml

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: minio-pv-claim
  labels:
    app: minio-storage-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```

Deploy MinIO on Kubernetes - Without Specific Operator

minio-deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: minio-deployment
spec:
  strategy:
    type: Recreate
  selector:
    matchLabels:
      app: minio
  template:
    metadata:
      labels:
        app: minio
```

Deploy MinIO on Kubernetes - Without Specific Operator

minio-deployment.yaml

```
spec:
  volumes:
    - name: storage
      persistentVolumeClaim:
        claimName: minio-pv-claim
  containers:
    - name: minio
      # Pulls the default MinIO image from Docker Hub
      image: minio/minio
      args:
        - server
        - /storage
```

Deploy MinIO on Kubernetes - Without Specific Operator

minio-deployment.yaml

```
env:
  # MinIO access key and secret key
  - name: MINIO_ACCESS_KEY
    value: "admin123"
  - name: MINIO_SECRET_KEY
    value: "admin123"
ports:
  - containerPort: 9000
  # Mount the volume into the pod
volumeMounts:
  - name: storage # must match the volume name,
    above
    mountPath: "/storage"
```

Deploy MinIO on Kubernetes - Without Specific Operator

minio-service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: minio-service
spec:
  type: NodePort
  ports:
    - port: 9000
      targetPort: 9000
      protocol: TCP
  selector:
    app: minio
```


Deploy MinIO on Kubernetes

With Specific Operator

minio-usage.sh

```
# install operator
chmod +x kubectl-minio_4.0.3_linux_amd64
cp kubectl-minio_4.0.3_linux_amd64
  /usr/local/bin/kubectl-minio

# init
kubectl minio init

# deployment
kubectl minio tenant create tenant1 --namespace
  minio-tenant-1 --storage-class longhorn --servers 2
  --volumes 2 --capacity 40Gi
```

Deploy MinIO on Kubernetes

With Specific Operator

minio-usage.sh

```
# install operator
chmod +x kubectl-minio_4.0.3_linux_amd64
cp kubectl-minio_4.0.3_linux_amd64
  /usr/local/bin/kubectl-minio

# init
kubectl minio init

# deployment
kubectl minio tenant create tenant1 --namespace
  minio-tenant-1 --storage-class longhorn --servers 2
  --volumes 2 --capacity 40Gi
```

Deploy MinIO on Kubernetes - With Specific Operator

minio-deployment.yaml

```
apiVersion: miniocontroller.min.io/v1beta1
kind: MinIOInstance
metadata:
  name: minio
spec:
  metadata:
    labels:
      app: minio
    annotations:
      prometheus.io/path: /minio/prometheus/metrics
      prometheus.io/port: "9000"
      prometheus.io/scrape: "true"
  image: minio/minio:RELEASE.2019-07-24T02-02-23Z
  credsSecret:
    name: minio-creds-secret
  replicas: 2
```

Deploy MinIO on Kubernetes - With Specific Operator

minio-deployment.yaml

```
requestAutoCert: false
certConfig:
  commonName: ""
  organizationName: []
  dnsNames: []
env:
  - name: MINIO_BROWSER
    value: "on"
resources:
  requests:
    memory: 512Mi
    cpu: 250m
liveness:
  httpGet:
    path: /minio/health/live
    port: 9000
```


Deploy MinIO on Kubernetes - With Specific Operator

minio-deployment.yaml

```
    initialDelaySeconds: 120
    periodSeconds: 20
volumeClaimTemplate:
  metadata:
    name: data
  spec:
    accessModes:
      - ReadWriteOnce
    resources:
      requests:
        storage: 10Gi
```

Operator Implementation

From Scratch or With Framework



Operator Implementation - minio operator v1.0.0

pkg/apis/miniocontroller/v1beta1/types.go

```
type MinIOInstanceSpec struct {  
    Image string `json:"image,omitempty"`  
    Replicas int32 `json:"replicas,omitempty"`  
    // ...  
    VolumeClaimTemplate *corev1.PersistentVolumeClaim  
    `json:"volumeClaimTemplate,omitempty"`  
    // ...  
    Mountpath string `json:"mountPath,omitempty"`  
    Subpath string `json:"subPath,omitempty"`  
    Liveness *corev1.Probe `json:"liveness,omitempty"`  
    Readiness *corev1.Probe `json:"readiness,omitempty"`  
    // ...  
}
```

Operator Implementation - minio operator v1.0.0

pkg/controller/cluster/controller.go

```
func (c *Controller) syncHandler(key string) error {
    // ...
    mi.EnsureDefaults()

    // ...
    if mi.Spec.Replicas != *ss.Spec.Replicas {
        glog.V(4).Infof("MinIOInstance %s replicas: %d,
        StatefulSet replicas: %d", name, mi.Spec.Replicas,
        *ss.Spec.Replicas)
        ss = statefulsets.NewForCluster(mi, svc.Name)
        _, err =
        c.kubeClientSet.AppsV1().StatefulSets(mi.Namespace).Update(ss)
    }
    // ...
}
```


Operator Implementation - minio operator v1.0.0

pkg/apis/miniocontroller/v1beta1/helper.go

```
func (mi *MinIOInstance) EnsureDefaults() *MinIOInstance {
    if mi.Spec.Replicas == 0 {
        mi.Spec.Replicas = constants.DefaultReplicas
    }
    if mi.Spec.Image == "" {
        mi.Spec.Image = constants.DefaultMinIOImage
    }
    if mi.Spec.Mountpath == "" {
        mi.Spec.Mountpath = constants.MinIOVolumeMountPath
    } else {
        mi.Spec.Mountpath = path.Clean(mi.Spec.Mountpath)
    }
    // ...
    return mi
}
```

Operator Implementation - minio operator v1.0.0

pkg/resources/statefulsets/statefulset.go

```
func NewForCluster(mi *miniov1beta1.MinIOInstance,
    serviceName string) {
    // ...
    if mi.Spec.VolumeClaimTemplate == nil {
        podVolumes = append(podVolumes,
            corev1.Volume{Name: constants.MinIOVolumeName,
                VolumeSource:
                    corev1.VolumeSource{EmptyDir:
                        &corev1.EmptyDirVolumeSource{Medium: ""}}})
    }
    // ...
    containers := []corev1.Container
    {minioServerContainer(mi, serviceName)}
    // ...
    return ss
}
```

Summay

Summary

So far, we know...

- the basic concept of kubernetes
- the interface and usage
- the principle of operator

Not yet introduced...

- kubernetes storage
- minio with kubernetes

Thanks for your listening!